

Note: Batch Size and Convergence

The loss update for mini-batch gradient descent follows the standard rule below.

```
theta = theta - eta * gradient(J(theta))
```

A reference implementation of the update step:

```
def sgd_step(theta, grad, eta=0.01):  
    return theta - eta * grad
```